Roarmic — Business Plan

Turn every phone in the room into a controlled microphone.

A SkillSimm Production Case · Plan date: July 2026 · Planning horizon: FY2026 Q3 – FY2029 Q2

---

1. Executive Summary

The problem. Rooms with more than about thirty people have a voice problem. One person has a microphone; everyone else is inaudible. The fixes available today are all bad: pass a handheld mic down a row, hire an AV operator, buy a $2,000–$8,000 wireless mic kit, or let people shout. In classrooms, town halls, churches, training rooms, and factory briefings, the result is the same — participation collapses to whoever is standing at the front.

The solution. Roarmic turns the phone already in every attendee's pocket into a controlled room microphone. Participants scan a QR code, join a session in the browser with no app install, and raise a hand. A moderator on the room computer decides who goes live, and that person's voice plays out through whatever speaker, amplifier, or projector audio the computer is already connected to. One live speaker at a time, with instant mute and stop-all — the control model is the product, not the audio capture.

Status. Roarmic is a working system, not a concept. The product runs today as a hosted web service with four screens — participant join, moderator control hub, projector display, and account login — and is live at a public address. Sessions, participants, and a complete log of moderation actions already persist. The remaining work to a commercial launch is reliability hardening, billing, and self-serve session provisioning — not core invention.

Market. The immediate serviceable market is the ~1.2M US classrooms, ~350K US congregations, and several hundred thousand corporate training and meeting rooms that have a computer wired to a speaker and no AV staff. At $99/room/year, capturing 0.5% of just the US classroom and congregation segments is roughly $7.5M of annual recurring revenue.

Business model. Self-serve SaaS. Free trial (10-minute sessions, 10 participants), Classroom at $9/month or $99/year, School/Team at $29/month or $299/year, and a $19 one-time Event Pass. A later Voice Intelligence tier — transcription, speaker attribution, and AI debrief tied to SkillSimm workflow steps — is the institutional upsell at $2,000–$8,000 per organization per year.

Financial highlights (modeled, assumptions stated in §9).

Measure	FY1	FY2	FY3
Paying rooms / orgs	420	2,600	8,300
Revenue	$41K	$268K	$1.04M
Gross margin	88%	90%	91%
Operating result	($64K)	($21K)	+$310K

FY1 ends June 2027; FY2 June 2028; FY3 June 2029.

The ask. $150,000 of seed or founder capital to reach cash-flow breakeven in FY3 Q1, allocated to one full-time engineer, connection reliability work, and $30K of segment-specific demand generation. The business is deliberately capital-light: infrastructure cost per paying room is under $1/year because live audio does not travel through our servers.

Why this wins. Every competitor sells either hardware (expensive, per-room, needs an operator) or text-only audience engagement (Slido, Mentimeter — questions get *typed*, not *spoken*). Roarmic is the only product that gives a moderator granular, low-friction control over *live voice* using equipment the

room already owns. And because Roarmic is a SkillSimm product, spoken input can become structured data attached to a workflow step — a capability no microphone can offer.

---

2. Products and Services

2.1 What the product does

Roarmic is a controlled voice-transfer system for a physical room. The path is deliberately short: a participant speaks into their phone, the audio travels across the room's own network to the computer at the front, and that computer plays it through the speaker, amplifier, or PA it is already connected to.

Our service coordinates the session — creating it, issuing the QR and short code, tracking who has joined, deciding who is live, carrying moderator commands, and enforcing plan limits. It does not carry the conversation. This is a deliberate choice with three consequences: delay stays short enough for natural turn-taking, bandwidth cost stays near zero as we scale, and we never become the custodian of a recording we did not need.

2.2 The four screens (all shipping today)

Screen	Who uses it	What it does
Participant join	Anyone with a phone	Scan, name, join, raise hand, speak when given the floor. No app, no account.
Control hub	Teacher or facilitator	Live roster, give and take the floor, silence, remove, lock the room, set time limits.
Projector view	The room	Big-screen "Now speaking", countdown, raised-hand queue. Read-only.
Account login	Account owner	Protected access to the control hub and projector view.

The control hub is the product's center of gravity. It is the device physically connected to the room's sound output, and it is where the moderation model lives: **exactly one participant is live by default.** That single constraint is what prevents the acoustic feedback, crosstalk, and chaos that make naive "phone as microphone" apps unusable in a real room.

2.3 Feature roadmap by release

v1.0 — Controlled Room Mic (launch scope, Q4 2026)

- QR or short-code join, no install and no account for participants
- Push-to-talk and a raise-hand queue
- Playback through the room computer's chosen speaker output
- Speaking indicator for each participant
- One live speaker; moderator switch, mute, unmute, end speaker, stop-all
- Room lock and per-speaker time limit
- A complete session and participation record
- Secure connections everywhere, so participant phones can use their microphones
- Billing, plan limits, and self-serve session creation

v1.5 — Reliability and Scale (H1 2027)

- A fallback connection path for guest networks that block device-to-device traffic

- Echo and noise handling tuned for playback in a live room
- Multi-room accounts, 5 admins, reusable session codes
- Event Pass purchase flow for one-off conferences and town halls
- Basic session report: who spoke, how long, how many hands raised

v2.0 — Voice Intelligence (H2 2027)

- Live transcription with speaker and team attribution
- Team-based routing ("Team 7 has the floor for 20 seconds")
- AI summary per speaker and per session
- Exportable debrief report
- SkillSimm integration: a spoken answer becomes the input to a workflow step, is evaluated, and appears in the instructor debrief dashboard

v2.5 — Operations Mode (2028)

- Emergency broadcast: one tap to all room speakers, logged
- Zone and multi-room broadcast for warehouses, campuses, care facilities
- Role permissions and single sign-on for institutional accounts

2.4 Intellectual property and defensibility

There is no patent position here and we do not claim one. The defensibility is:

1. **The moderation model.** The rules that make a 180-person room usable — one live speaker, a hand queue, time limits, an instant kill switch — are earned through field iteration, not invented in a specification.
2. **Reliability on hostile networks.** School and church Wi-Fi routinely blocks devices from reaching each other, hides behind sign-in pages, and filters aggressively. Making the join work anyway is unglamorous, difficult, and a real barrier.
3. **The SkillSimm data path.** Competitors can amplify a voice. They cannot turn that voice into a scored, attributed, auditable decision record attached to a workflow step. That path is proprietary to the SkillSimm platform.

---

3. Market Analysis

3.1 Industry context

Roarmic sits at the intersection of three markets, and is priced as none of them.

Commercial AV. Room audio is a hardware business: wireless mic systems, ceiling arrays, mixers, and installation labor. A single wireless handheld or lavalier system for a classroom runs $400–$1,200; a moderated audience Q&A system for an auditorium runs $3,000–$15,000 installed. The market is large and growing, but the purchase requires capital budget, procurement, and usually an integrator.

Audience engagement software. Slido, Mentimeter, Poll Everywhere, and Wooclap have proven that attendees will pull out a phone and participate through a browser during a live session. Slido alone reports usage across millions of events. Critically, **all of them route participation through text.** A question is typed, upvoted, and read aloud by the moderator. The audience's own voice never reaches the room.

Conferencing. Zoom, Teams, and Meet solved remote voice completely and in-room voice barely at all. "Bring your own device" room modes exist but are designed to prevent feedback by muting local devices — the opposite of what an in-room speaker needs.

Roarmic's opening is the gap between them: the AV answer is too expensive and needs an operator; the engagement answer is silent; the conferencing answer is designed for remote participants. Nobody sells cheap, controlled, in-room *voice*.

3.2 The problem, stated precisely

Five conditions predict acute pain. Where three or more are present, Roarmic goes from convenience to necessity:

1. **Many people, one microphone.** Passing a handheld through a lecture hall costs 30–60 seconds per question and suppresses participation from the middle of the room.
2. **Rotating speakers, not one presenter.** Traditional AV assumes a presenter. Group report-outs, panel Q&A, team decisions, and open forums assume the opposite.
3. **No AV staff.** Most schools, churches, training centers, and small conference venues have nobody whose job is sound. Setup friction kills adoption of anything harder than a QR code.
4. **No capital budget.** A teacher, volunteer coordinator, or workshop facilitator can approve $9/month personally. They cannot approve $1,500 without a procurement cycle.
5. **The voice needs to become a record.** Compliance training, safety briefings, public meetings, and graded participation all need evidence that something was said, by whom.

3.3 Market sizing (US, with a stated bottom-up method)

We size by *rooms with a computer connected to a speaker and no AV operator*, because that is the exact precondition for Roarmic to work with zero added hardware.

TAM — total addressable rooms (US):

Segment	Rooms	Basis
K-12 and higher-ed classrooms	1,200,000	~3.8M US teachers, discounted for shared rooms
Congregations	350,000	US congregations, most volunteer-run AV
Corporate training and meeting rooms	500,000	Firms of 50–1,000 with training space
Event venues, halls, civic chambers	100,000	Municipal and commercial meeting space
Total	**~2,150,000**	

At the $99/year Classroom price, **TAM ≈ $213M/year** on the basic tier alone.

SAM — realistically reachable in the plan horizon. Restrict to rooms where (a) the room computer is present and connected, (b) the buyer can self-serve, and (c) the network permits local audio. We assume 35% of TAM clears all three: **~750,000 rooms ≈ $74M/year.**

SOM — 3-year target. 8,300 paying rooms/orgs by FY3 ≈ **1.1% of SAM ≈ $1.04M ARR.** This is deliberately conservative: it assumes no enterprise channel and no district-level bulk purchase, both of which are upside.

Voice Intelligence expansion. The v2.0 institutional tier at $2,000–$8,000/org/year targets a much smaller, much higher-value population — roughly 25,000 US institutions running graded participation, compliance training, or public comment. At $3,500 average, that is an additional **$87M TAM** which we model at only 40 accounts in FY3.

3.4 Target segments, ranked by go-to-market priority

Priority 1 — Large-enrollment university courses. 180 students, one instructor, no TA, participation is part of the grade. Highest pain, fastest word-of-mouth, and the segment where the v2.0 attribution and debrief features convert an individual subscription into a department purchase. Buyer: the instructor. Budget: personal or small departmental. Cycle: days.

Priority 2 — Churches and community congregations. Volunteer-run AV, weekly recurring use, budget-sensitive, and unusually strong peer referral networks between congregations. Buyer: AV volunteer or administrator. Cycle: one meeting.

Priority 3 — Corporate training and workshop facilitators. Team report-outs are the canonical use case: "every group has 30 seconds to report." Facilitators run the same format across many client rooms, so one customer generates many sessions. Willing to pay more than $9. Buyer: L&D lead or independent facilitator.

Priority 4 — Public meetings and town halls. Moderated public comment with a queue and a record is a direct fit for the raise-hand and participation-log design. Longer sales cycle, procurement involved, but very high retention once installed.

Priority 5 — Operations environments (factory briefings, warehouses, construction, senior care). Highest strategic value in v2.5 with emergency broadcast mode, but requires reliability and support commitments we should not make in FY1.

3.5 Trends supporting the thesis

- **Phones are universal in the audience.** The bring-your-own-device assumption that made Slido viable in 2016 is now unremarkable in classrooms, churches, and factories alike.
- **Browsers can do this now.** Live audio in a web page works on every current phone with no app store, no install, and no IT deployment.
- **AV budgets are flat while room count grows.** Hybrid work multiplied meeting spaces without multiplying AV staff.
- **Participation is increasingly measured.** Accreditation, compliance training, and public-comment requirements all push toward evidence of who spoke.

---

4. Competition

4.1 Competitive landscape

Competitor	Price	Where they beat us	Where we beat them
Wireless mic systems (Shure, Sennheiser)	$400–$8,000 per room	Audio quality, no network needed	Cost, setup time, no moderation, no record
Catchbox (throwable mic)	$600–$1,000	Delightful, needs no network	One mic, no queue, no data, no control
Ceiling array systems	$2,000–$15,000 installed	Hands-free, professional	Needs an integrator; no per-speaker control
Slido, Mentimeter, Poll Everywhere	$10–$100/mo	Brand, polls, huge install base	**Text only.** The audience never speaks aloud
Zoom / Teams room modes	Bundled	Remote participants, enterprise trust	Built to *suppress* in-room device audio
Consumer "phone as mic" apps	Free–$5	Cheap	One user, no moderation, no record
Doing nothing (pass the mic)	$0	Free, familiar	Slow, exclusionary, no record

4.2 The honest competitive read

Our most dangerous competitor is inertia, not another product. Most rooms tolerate bad audio because the alternative is a purchase order. Our positioning must therefore beat "do nothing" on setup time, not beat Shure on fidelity. If a teacher cannot go from signup to a student speaking through the classroom speaker in under five minutes, we lose to inertia every time.

Our most dangerous strategic threat is Slido or Mentimeter adding live voice. They have the audience-side distribution and brand permission to do it. Our defenses are: (a) speed — ship the moderation model and reliability work before they decide it is worth doing; (b) the room-audio output path, which is genuinely fiddly work they have no reason to have built; and (c) the SkillSimm data path, which is not an audio feature at all.

4.3 Sustainable advantages

1. **Zero added hardware.** We use the amplifier the room already owns. Every hardware competitor's price floor is our entire annual subscription.
2. **Moderation-first design.** The value is in who does *not* get to speak. One live speaker, a hand queue, an instant kill switch, time limits, room lock.
3. **Zero-install participation.** Scan, name, in. No app store review, no IT deployment, no account for participants.
4. **Marginal cost near zero.** Live audio does not pass through our infrastructure, so gross margin holds above 88% at any scale (see §9.3).
5. **The SkillSimm bridge.** A spoken answer becomes a transcript, attributed to a team, attached to a workflow step, evaluated, and surfaced in a debrief. This turns a $99 AV utility into a $3,500 institutional assessment tool, and no AV or engagement vendor is positioned to follow.

---

5. Marketing and Sales

5.1 Positioning

We deliberately do **not** lead with "use your phone as a microphone." That reads as a convenience app and invites comparison to free utilities.

Primary positioning: *"Instant voice routing for groups, teams, and live rooms."*

Segment message (classroom/congregation): "Turn every phone in the room into a controlled microphone. No mic kit, no passing a mic, no AV person — use the computer already connected to your speaker."

Platform positioning (SkillSimm, v2.0+): *"Voice-to-workflow infrastructure for live human–AI simulations."*Here the voice is not the product; it is the input layer for coordination, assessment, and decision records.

5.2 Pricing

Priced as a room-audio utility, not a classroom platform. The buyer is paying for one outcome, so the price must be approvable without a meeting.

Plan	Price	For	Includes
Free Trial	$0	Testing one room	10-minute sessions, 10 participants, 1 admin
Classroom	$9/mo or $99/yr	One teacher or facilitator	Unlimited sessions, full moderation, 1 admin
School / Team	$29/mo or $299/yr	Multiple rooms or instructors	5 admins, shared account, reusable codes
Event Pass	$19 one-time	Conferences and town halls	One session up to 8 hours, full moderation
Voice Intelligence	$2,000–$8,000/yr	Institutions (2027)	Transcription, attribution, AI debrief, exports

Launch simplification. Ship only Free Trial and Classroom. Add School/Team only after several individual customers ask for multiple admins, and Event Pass only after inbound event demand appears. Every plan added before it is demanded is a conversion tax.

What we deliberately do not price separately in v1: polls, transcription, AI summaries, recording, templates, workflows, hardware. Those are the product's future, not its paywall — and unbundling them at launch makes the offer hard to explain.

Pricing rationale. $99/year is below the discretionary-approval threshold for a teacher, a church volunteer, or a workshop facilitator, and it is ~10% of the cheapest wireless mic system it replaces. We are pricing to remove the decision, not to maximize ARPU. ARPU rises later through seat expansion and the Voice Intelligence tier, not through raising the base price.

5.3 Customer acquisition

Community-led (primary, FY1). Roarmic solves a problem that people already complain about publicly. Genuine participation — not link-dropping — in:

- AV professional forums, where "cheap classroom mic?" is a recurring question
- Education technology and instructor communities — the primary buyer
- Audio and sound-engineering communities — credibility and technical scrutiny
- Church tech communities and volunteer AV groups
- Facilitator and corporate learning-and-development communities

Target: 40 substantive community interactions per month, 15% resulting in a trial.

Content and SEO (compounding, from Q4 2026). We write for the exact query the buyer types: *"cheap classroom microphone for large lecture hall"*, *"audience Q&A without mic runners"*, *"church wireless mic alternative"*, *"how to let students speak through classroom speakers"*. Target 24 articles in 12 months, each mapping to one segment pain.

Product-led growth. The free trial is the demo. The QR code is the viral surface: every session puts our join page in front of 20–200 phones, and the join page shows what the product is. A participant who becomes a facilitator elsewhere is our lowest-cost acquisition channel. Target 12% of new trials from participant-to-host conversion by FY2.

Video demonstration. This product is almost impossible to explain in text and trivially obvious in a 45-second video: room, QR code, phone, voice out of the ceiling speaker. Short-form video is the highest-leverage channel and receives 30% of the FY1 marketing budget.

Institutional pilots (FY2+). For the Voice Intelligence tier, run 3–5 paid pilots with university teaching-and-learning centers. A published pilot result ("participation rose from 6 speakers to 41 per session") is the asset that opens department-wide purchases.

5.4 Sales process

FY1 is **entirely self-serve**: trial, in-product upgrade, card payment. No sales calls, because a $99 average contract cannot support one. The only human-touch motion is inbound support during a live session, which we treat as a sales function — a failed first session is a lost customer and a bad word-of-mouth event.

FY2 adds a lightweight founder-led motion for School/Team and Voice Intelligence: 30-minute pilot scoping, a single-page quote, and a semester-long pilot terminating in an invoice.

5.5 Retention

Retention is the whole model at this price. Two mechanics:

- **Habit.** A teacher who uses Roarmic twice a week for a semester will renew. First-30-day session count is our leading retention indicator; target 4+.
- **Seasonality.** Education churns at semester boundaries. We counter with annual billing (priced 8% below monthly to push toward it) and a "your room is ready" re-activation sequence timed to the academic calendar.

Target logo retention: 68% annual in FY1, 75% by FY3 as reliability improves.

---

6. Operations

6.1 How the product is delivered

Roarmic is a hosted web service. Customers reach it at a web address, participants reach it by scanning a code, and there is nothing to install on either side. The service runs on shared SkillSimm infrastructure alongside the rest of the portfolio, which keeps per-product operating cost near zero and means a single operations discipline covers every product we ship.

The economics follow from one design choice: our servers coordinate sessions but do not carry the audio. Adding a thousand rooms adds coordination traffic measured in kilobytes, not conversation traffic measured in gigabytes.

6.2 The first constraint: microphone permission

Phone browsers grant microphone access only over a secure connection. This is the single most important operational fact about the product, and it dictates the roadmap: until every customer session runs over a secure hosted address, participants cannot speak from their own phones — which is the core promise. The hosted service already meets this requirement; the v1.0 work is making every customer's session flow run entirely through it, so no customer is ever asked to change a setting on their own machine.

6.3 The second constraint: room networks

School, church, and conference Wi-Fi frequently prevents devices on the same network from reaching each other, and sign-in pages and firewalls compound the problem. Mitigations, in order of implementation:

1. Detect the condition at join time and tell the user plainly what is wrong (v1.0)
2. Fall back to a relayed connection when the direct one fails (v1.5)
3. Publish a one-page network requirements document for IT departments (v1.0)
4. Offer a mode where the room computer serves the session directly (v1.5)

We budget relay bandwidth in the financial model (§9.3) assuming 20% of sessions require it — a deliberately pessimistic assumption.

6.4 The third constraint: acoustics

A phone microphone in the same room as the loudspeaker playing it is a feedback loop. The one-live-speaker rule is the primary defense; the browser's own echo handling is the second. Field guidance — speaker not pointed at the audience's phones, participant holds the phone near the mouth, a ceiling on hub volume — ships as in-product tips in v1.0, with tuned echo and noise handling in v1.5.

6.5 Release practice

Releases ship weekly during FY1 and never during North American school hours (08:00–16:00 local). A live session interrupted by an update is a churn event, and the release calendar treats it as one.

6.6 Support

In-app help and a documented troubleshooting page for the three failure modes that account for nearly all incidents: microphone permission denied, a network that blocks the connection, and no audio output selected on the room computer. Email support with a one-business-day target on paid plans, plus a "live session emergency" address monitored during business hours. Support load is the primary constraint on how fast we can grow in FY1, which is why v1.5′s reliability work precedes FY2′s growth spend.

6.7 Data, privacy, and compliance

- **We do not record audio in v1.** This is both a product position and a compliance simplification, and it is a selling point to schools.
- Stored data is minimal: session details, a display name per participant, and a log of moderation actions. No accounts and no personal information for participants.
- Because customers are schools, student-privacy expectations apply from day one: documented retention (12 months, then purge), no third-party data sharing, and a data-processing agreement available on request.
- When transcription arrives in v2.0, it is **opt-in per session**, visibly indicated to participants on the join screen, and covered by explicit consent language. Recording consent law varies by state; the product must show participants when they are being transcribed, without exception.

---

7. Milestones and Metrics

7.1 Milestones

Phase	Window	Deliverables	Exit criteria
P0 — Harden	Q3 2026	Secure participant microphone access; hosted session provisioning; network diagnostics	10 consecutive real sessions, zero audio failures
P1 — Launch	Q4 2026	Billing; trial limits; self-serve signup; onboarding; landing page; demo video	50 paying rooms; <5% first-session failure rate
P2 — Reliability	H1 2027	Relay fallback; echo and noise tuning; School/Team; Event Pass; session report	420 paying rooms; 68% retention; NPS ≥ 40
P3 — Voice Intelligence	H2 2027	Transcription; attribution; team routing; AI summary; SkillSimm workflow hook	3 paid pilots; first $2K+ contract signed
P4 — Institutional	2028	Single sign-on; role permissions; zone broadcast; campus licensing	2,600 paying rooms; 12 Voice Intelligence accounts
P5 — Expansion	FY3	Operations vertical (factory, care, construction); partner channel	$1M ARR; cash-flow positive

7.2 Metrics we manage to

Activation (the FY1 obsession)

- Signup to first successful session, target **under 10 minutes**
- Share of first sessions with at least one participant speaking, target **over 90%**
- Median time from scan to joined, target **under 20 seconds**

Engagement

- Sessions per paying room per month, target **8**
- Participants per session, median target **22**
- Speakers per session (the participation metric we sell), target **6+**

Revenue

- Trial to paid conversion, target **18%**
- Annual vs monthly plan mix, target **55% annual** by FY2
- ARPU, target $92 (FY1) rising to $126 (FY3) via plan mix, not price increases

Retention

- Logo retention, 68% (FY1) to 75% (FY3)

- Net revenue retention including School/Team upgrades, target **95%+** by FY3
- Sessions in the first 30 days (leading indicator), target **4+**

Reliability (the metric that gates growth spend)

- Session audio failure rate, target **under 3%**
- Joins blocked by the room network, target **under 5%** after relay fallback
- Uptime during scheduled sessions, target **99.5%**

---

8. Company Overview and Team

8.1 Structure

Roarmic operates as a product line of **SkillSimm**, sharing infrastructure, billing entity, and brand. It is deliberately *not* a separate company: the strategic value of Roarmic is partly standalone SaaS revenue and partly its role as the voice input layer for SkillSimm's simulation and workflow products. If Roarmic reaches $1M ARR on its own, spinning it into a subsidiary becomes a live option; before then, separation would only duplicate overhead.

Legal form: LLC, with conversion to C-corp contemplated only if institutional investment is raised.

8.2 History

- **Q1–Q2 2026** — Problem identified in classroom settings; product design, pricing strategy, demand-channel research, and positioning thesis documented.
- **Q2 2026** — Working system built: all four screens, session and participation records, and branded session codes.
- **Q3 2026** — Deployed to production on shared SkillSimm infrastructure. Currently in pre-commercial field testing.

8.3 Team

Founder / Product and Engineering (full-time). Owns product definition, build, deployment, and early customer development. Has built and shipped the current system end to end.

Planned hires:

Role	Timing	Rationale	Cost
Full-stack engineer	Q4 2026	Connection reliability — the work that gates growth	$95K
Customer success (half-time)	Q3 2027	Support load is the FY2 growth constraint	$32K
Growth marketer (contract)	Q4 2026	Content, video, community — 20 hrs/week	$36K
Transcription engineer (contract)	Q3 2027	v2.0 Voice Intelligence	$45K (6 mo)

Advisors sought: one commercial AV integrator (to pressure-test where we lose to hardware), one university teaching-and-learning center director (institutional buying), and one K-12 IT administrator (network reality).

8.4 Operating principles

- **The room is the customer, not the org chart.** Optimize for a facilitator alone in a room five minutes before class, not for a procurement committee.

- **Moderation before capture.** Any feature that increases who can speak without increasing the moderator's control makes the product worse.
- **Don't own the audio.** Keep live voice off our servers. It is cheaper, faster, and a better privacy posture.
- **Reliability is marketing.** One failed session in front of 180 people costs more than any campaign buys.
- **Sell participation, not audio.** The purchase is controlled participation, faster coordination, and usable records.

---

9. Financial Plan and Forecasts

All figures are **modeled projections**, built bottom-up from the assumptions stated below. Fiscal years run July–June to match the academic calendar. FY1 = Jul 2026 – Jun 2027.

9.1 Startup requirements

Item	Amount	Notes
Engineering (1 FTE, 9 months)	$71,000	Reliability, billing, provisioning
Growth marketing contractor	$27,000	Content, video, community
Infrastructure (FY1)	$4,200	Hosting, data, relay, monitoring
Software and services	$3,600	Payments, error tracking, analytics, email
Paid demand generation	$18,000	Video production, targeted social, tests
Legal, accounting, privacy review	$9,000	Includes education data-handling review
Contingency (12%)	$15,800	
Total FY1 requirement	**$148,600**	Rounded ask: **$150,000**

9.2 Revenue model and assumptions

Stated assumptions:

- Blended ARPU begins at $92/year (mostly Classroom annual, some monthly churn-adjusted) and rises to $126/year by FY3 as School/Team grows to 22% of accounts.
- Trial to paid conversion of 18%, based on the free tier being genuinely useful but time-limited (10-minute sessions).
- Logo retention of 68% (FY1), 72% (FY2), 75% (FY3).
- Event Pass contributes 6% of FY2 revenue and 8% of FY3.
- Voice Intelligence launches H2 2027 at an average $3,400 contract; 4 accounts in FY2, 40 in FY3.
- No paid channel scaling until session failure rate is under 3%.

Revenue forecast:

Line	FY1	FY2	FY3
New paying rooms added	420	2,470	6,940
Retained from prior year	—	286	1,872
Ending paying rooms	**420**	**2,600**	**8,300**
Blended ARPU	$92	$99	$126
Subscription revenue	$38,600	$242,600	$960,000
Event Pass revenue	$2,400	$12,000	$79,000
Voice Intelligence revenue	$0	$13,600	$136,000
Total revenue	**$41,000**	**$268,200**	**$1,175,000**

(FY3 total revenue of $1.175M vs. the $1.04M headline in §1: the headline figure excludes Voice Intelligence, which is the riskiest line. Both are shown deliberately.)

9.3 Cost of revenue and gross margin

The defining economic fact: **live audio does not pass through our servers**, so cost per paying room is dominated by fixed hosting rather than usage.

Cost driver	Per room / year	Basis
Hosting and data	$0.55	Shared infrastructure, ~1,500 rooms per unit
Relay bandwidth	$0.90	20% of sessions relayed, 8 sessions/mo, 45 min avg
Payment processing	$3.60	~2.9% + $0.30 on a $99 charge
Support (allocated)	$5.80 to $3.10	FY1 to FY3; falls as reliability improves
Transcription (v2.0 only)	~$110 per account	Usage-based, institutional volumes
Blended cost	**~$11 to $11.50**	FY1 to FY3

Line	FY1	FY2	FY3
Revenue	$41,000	$268,200	$1,175,000
Cost of revenue	$4,900	$26,800	$105,800
Gross profit	**$36,100**	**$241,400**	**$1,069,200**
Gross margin	**88%**	**90%**	**91%**

9.4 Profit and loss forecast

Line	FY1	FY2	FY3
Revenue	$41,000	$268,200	$1,175,000
Cost of revenue	(4,900)	(26,800)	(105,800)
Gross profit	**36,100**	**241,400**	**1,069,200**
Engineering	(71,000)	(128,000)	(245,000)
Customer success	0	(32,000)	(78,000)
Marketing and demand gen	(45,000)	(72,000)	(310,000)
Infrastructure and tooling	(7,800)	(12,000)	(26,000)
G&A, legal, accounting	(9,000)	(18,000)	(42,000)
Contingency	(15,800)	0	(58,000)
Total operating expense	**(148,600)**	**(262,000)**	**(759,000)**
Operating result	**($112,500)**	**($20,600)**	**+$310,200**

(FY1 operating result is stated here on a full-cost basis; the $(64K) figure in §1 reflects the founder salary being deferred, which is the base case for FY1.)

9.5 Cash flow and breakeven

- **FY1:** Cash burn of ~$78K against $150K raised (founder compensation deferred). Ending cash ~$72K.
- **FY2:** Near-breakeven on operations; annual prepayments are a favorable working-capital dynamic — 55% of subscription revenue collected upfront means cash collections run ahead of recognized revenue. Ending cash ~$55K without additional raise.
- **FY3 Q1: Cash-flow breakeven**, at approximately 3,100 paying rooms.
- FY3 generates ~$310K of operating cash, which funds the operations vertical without a second raise.

Breakeven arithmetic (steady-state FY2 cost base of ~$262K): at $99 gross ARPU less $11 of cost, each room contributes $88/year. Breakeven ≈ $262,000 / $88 ≈ **2,980 paying rooms.**

9.6 Unit economics

Metric	FY1	FY3	Method
CAC	$58	$41	Marketing spend ÷ new paying rooms
ARPU	$92	$126	Blended across plans
Gross margin per room	$81	$115	ARPU less allocated cost
Average customer life	3.1 yrs	4.0 yrs	1 / (1 − retention)
LTV	**$251**	**$460**	Gross-margin basis
LTV : CAC	**4.3×**	**11.2×**	
CAC payback	8.6 months	4.3 months	

The model is viable specifically because CAC is low, and CAC is low specifically because the product is community-discovered and demo-viral rather than ad-purchased. **If CAC exceeds $90, the $99 price**

point stops working and the plan must shift to School/Team as the primary entry product. That is the single most important sensitivity in this plan.

9.7 Sensitivity analysis

Scenario	Change	FY3 revenue	FY3 result
Base	As modeled	$1,175,000	+$310,000
Downside	Conversion 12%, retention 60%, no v2.0	$498,000	($120,000)
Upside	Campus channel deal, 70 v2.0 accounts	$1,760,000	+$690,000
CAC stress	CAC rises to $95	$1,175,000	+$60,000
Reliability failure	Failure rate >8%, retention 50%	$610,000	($195,000)

9.8 Funding requirements and use of funds

Seeking $150,000 as founder capital, a small seed, or revenue-based financing.

Use	Amount	%
Engineering (reliability, billing, connection quality)	$71,000	47%
Demand generation (video, content, community)	$45,000	30%
Infrastructure and tooling	$12,000	8%
Legal, accounting, privacy compliance	$9,000	6%
Contingency	$13,000	9%

Milestone this capital buys: 420 paying rooms, session failure rate under 3%, and a validated $99 self-serve motion — which is the evidence required either to raise a proper seed for the institutional tier or to grow from operating cash.

9.9 Exit considerations

Realistic acquirers, in order of likelihood:

1. **Audience engagement platforms** (Slido/Cisco, Mentimeter, Poll Everywhere) — buying the live-voice capability they lack, at a strategic multiple on ARR.
2. **EdTech platforms** (Instructure, Echo360, Top Hat) — buying participation capture and the attribution data that follows.
3. **AV manufacturers** (Shure, Sennheiser, Biamp) — buying a software beachhead in rooms too small for their hardware.
4. **Independent operation.** At 90% gross margin and low capital intensity, a $2–4M ARR Roarmic is a durable cash-generating business with no obligation to sell.

---

10. Appendix

A. The participation record

Every session records who joined and left, who spoke and for how long, who raised a hand, and every moderation action a host took — locks, removals, and time limits.

This record is what makes participation reporting, the v2.0 debrief, and any future assessment product possible. It is already collected in production, which means the data asset begins accumulating before the feature that monetizes it ships.

B. Risk register

Risk	Likelihood	Impact	Mitigation
Room network blocks phone-to-computer audio	High	High	Detection at join; relay fallback; IT requirements doc
Acoustic feedback in real rooms	Medium	High	One live speaker; echo handling; placement guidance
Users hit microphone permission errors	High	Medium	Fully hosted secure flow; never expose workarounds
Slido or Mentimeter ships live voice	Medium	High	Speed; the room-audio path is hard; SkillSimm data moat
$99 price cannot support acquisition cost	Medium	High	Community-first GTM; shift entry plan if CAC > $90
Education seasonality creates churn cliffs	High	Medium	Annual billing; calendar reactivation; other segments
Support load caps growth	Medium	Medium	Reliability before growth spend; self-serve help; CS hire
Privacy concerns block school adoption	Medium	High	No recording in v1; opt-in transcription; documented retention
Single-founder key-person risk	High	High	Engineer hire Q4 2026; documented process; shared platform
Audio quality judged against hardware mics	Medium	Medium	Position against “no mic at all”; publish honest expectations

C. Key assumptions requiring validation

1. A facilitator will complete signup to first successful session in under 10 minutes.
2. Participants will scan a code and speak without prompting or embarrassment.
3. Direct phone-to-computer audio succeeds on at least 80% of real institution networks.

4. Trial to paid conversion reaches 18% at a 10-minute session limit.
5. Word of mouth within education and church communities produces CAC under $60.
6. Institutions will pay $2,000+ for participation attribution and AI debrief.

Assumptions 1–3 are testable within 60 days of P0 completion and should gate all growth spend. Assumption 6 gates the entire Voice Intelligence line and should be validated with three letters of intent before any transcription work begins.